

Introduction to R

732A99/TDDE01, Lecture 1c

What is R?



- Programming language for statistical computing and graphics
 - Developed in 1993 by Ross Ihaka and Robert Gentleman at University of Auckland
 - Managed by R Foundation and open-source
- Constantly increasing amount of software packages
 - Used in many fields and promoted by research

Software: RStudio

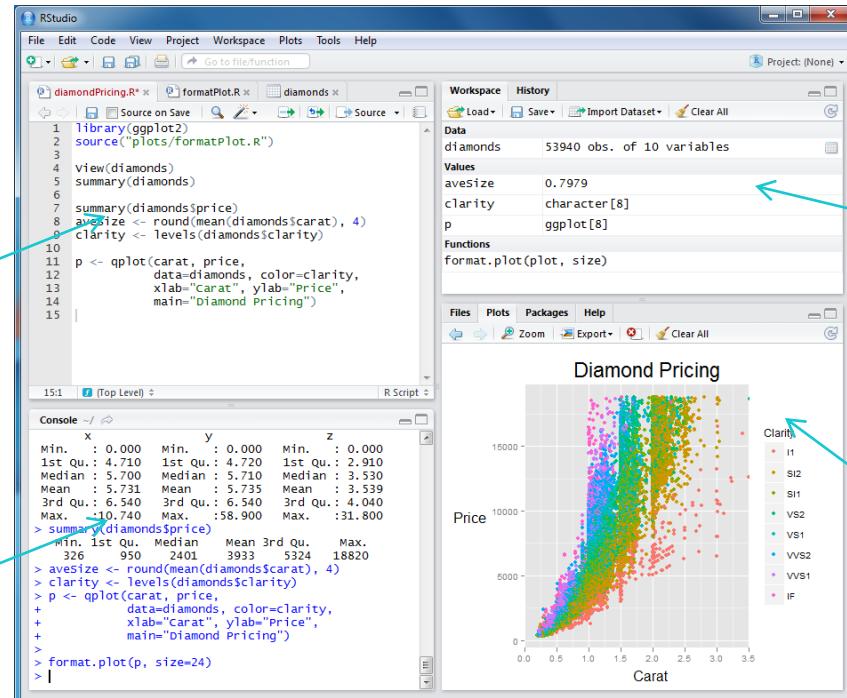
- Install R: <http://www.r-project.org/>
- Install RStudio: <http://rstudio.org/>

Program

Execution
console

Workspace

Plots



RStudio: Basics

- Important to know:
 - Create a new file and save it (File menu)
 - Running one line or entire code (Edit menu)
 - Running one line in console
 - Workspace (Observe, Save, Clear)
 - Setting current directory (Tools)
 - Installing new package (Packages tabs)

RStudio: Basics

- Specific function
 - `help(function)`
- Help browser
 - `help.start()`
- Search for something in help
 - `help.search("expression")`
- Quick reminder of function arguments:
 - `args(function)`
- Examples of how to use function:
 - `example(function)`
- If some method is not installed on the computer:
 - `RSiteSearch("expression")`

Getting started in R

- Each line of code forms a separate command
- Variables are initialized with -> or <- or =, e.g.
 - a <- 3
 - 3 -> a
 - a = 3
- R is case-sensitive (A and a!)
- Comments start with #, e.g.
 - # R is a great programming language!

Vectors

- Creating a vector: `x <- c(1, 3)`
- View the content of a vector: `x` or `print(x)`
- Create an empty vector: `Y <- numeric(10)`

```
> x<-c(1,3)
> x
[1] 1 3
> print(x)
[1] 1 3
>
```

```
> Y=numeric(10)
> Y
[1] 0 0 0 0 0 0 0 0 0 0
```

Sequences

- Sequences can be defined either by
 - `start:end`
 - `seq(from, to, by)`

R R Console

```
> f<-3:5
> f
[1] 3 4 5
> g<-seq(from=3, to=7, by=0.5)
> g
[1] 3.0 3.5 4.0 4.5 5.0 5.5 6.0 6.5 7.0
> |
```

Vector Operations

- Indexing: [] (Start at 1!)
- Element-wise operations:
 - +-*/^
 - Log, exp, sin, cos ...
- Vector-wise operations:
 - Number of elements: length()
 - Sum of all elements: sum()
 - Max(), min(), which.min(), which.max()
 - Sort(), Order()
 -

```
> a=1:5  
> b=c(1,4,-1,3,0)  
  
> b[1]  
[1] 1  
> b[2:4]  
[1] 4 -1 3  
> b[-2]  
[1] 1 -1 3 0  
  
> a+b  
[1] 2 6 2 7 5  
> a*b  
[1] 1 8 -3 12 0  
> b+4  
[1] 5 8 3 7 4  
> length(a)  
[1] 5  
> sum(a^2)  
[1] 55  
> max(b)  
[1] 4  
> which.max(b)  
[1] 2  
> order(b)  
[1] 3 5 1 4 2  
> sort(b)  
[1] -1 0 1 3 4
```

Vector operations

- Replication : `rep(vector, times)`

```
> v1=rep(3,5)
> v1
[1] 3 3 3 3 3
> v2=rep(c(3,4),2)
> v2
[1] 3 4 3 4
```

Matrices

- Create a matrix:

```
a <- matrix(values, nrow=n, ncol=c)
```

→ "values" should be column-wise

→ if empty matrix, simply replace by 0

R R Console

```
> a<-matrix(c(2,1,1,-1),nrow=2,ncol=2)
> a
 [,1] [,2]
[1,]    2    1
[2,]    1   -1
> |
```

```
> m=matrix(0,nrow=2,ncol=3)
> m
 [,1] [,2] [,3]
[1,]    0    0    0
[2,]    0    0    0
> |
```

Matrix operations

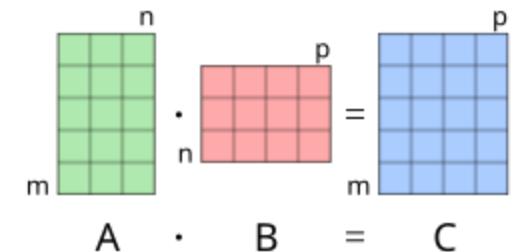
- Indexing for matrices:
 - Positive index: e.g., `x[1, 6]` or `x[2:10,]`
 - Negative index: e.g., `x[2, -(1:5)]` (row 2 and all columns except 1:5)
 - Entire column or row: `y=x[2,]`
 - Extraction: `x[x>5]` (all values above 5)

Matrix operations

Usual vector operations can be applied, as well as matrix-specific ones:

```
> m4=matrix(c(1,2,0,1), nrow=2)
> m5=matrix(c(2,2,5,1), nrow=2)
> m4
     [,1] [,2]
[1,]    1    0
[2,]    2    1
> m5
     [,1] [,2]
[1,]    2    5
[2,]    2    1
> m4*m5
     [,1] [,2]
[1,]    2    0
[2,]    4    1
```

```
> x<-c(1,2)
> a<-matrix(c(2,1,1,-1),2,2)
> b<-matrix(c(1,0,1,1),2,2)
> y=a%*%x
> y
     [,1]
[1,]    4
[2,]   -1
> c=a%*%b
> c
     [,1] [,2]
[1,]    2    3
[2,]    1    0
> |
```



Matrix operations

- Matrix operators/functions:
 - Transpose $b = a^T \rightarrow b=t(a)$
 - Inverse $b = a^{-1} \rightarrow b=solve(a)$
 - Solve $d = a^{-1}b \rightarrow d=solve(a,b)$

```
> a
      [,1] [,2]
[1,]    2    1
[2,]    1   -1
> t(a)
      [,1] [,2]
[1,]    2    1
[2,]    1   -1
> solve(a)
      [,1]      [,2]
[1,] 0.3333333  0.3333333
[2,] 0.3333333 -0.6666667
> |
```

Matrix operations

- Get Dimensions:

- `dim(mat)`
- `nrow(mat)`
- `ncol(mat)`

```
> m2
 [,1] [,2]
[1,] 3 3
[2,] 4 4
[3,] 3 3
[4,] 4 4
> ns=dim(m2)
> ns[2]
[1] 2
```

- Row/column statistics:

- `colMeans()`
- `rowMeans()`
- `colSums()`
- `rowSums()`

```
> m4=matrix(c(1,3,5,2,4,6),nrow=3)
> m4
 [,1] [,2]
[1,] 1 2
[2,] 3 4
[3,] 5 6
> colMeans(m4)
[1] 3 4
> rowSums(m4)
[1] 3 7 11
```

- Apply a function over vector/matrix: `sapply()`, `apply()`, `lapply()`

→ Normally used when function works only element-wise

```
> sapply(v2,log)
[1] 1.098612 1.386294 1.098612
> log(v2)
[1] 1.098612 1.386294 1.098612
```

Matrix operations

- Create confusion matrix (classification):
`table(X, Y)`
- Extract diagonal:
`diag(X)`

```
> X=c(1,3,1,1,2,3,1,2,2,2,1,1,3)
> xfit=c(2,3,2,1,2,3,1,2,2,2,1,1,1)
> t1=table(xfit,X)
> t1
      X
xfit 1 2 3
      1 4 0 1
      2 2 4 0
      3 0 0 2
> t1[1,1]
[1] 4
> diag(t1)
1 2 3
4 4 2
>
```

Matrix operations

- Replication: `matrix()`

```
> m1=matrix(1,nrow=2,ncol=2)
> m1
     [,1] [,2]
[1,]    1    1
[2,]    1    1
> m2=matrix(v2,nrow=4,ncol=2)
> m2
     [,1] [,2]
[1,]    3    3
[2,]    4    4
[3,]    3    3
[4,]    4    4
> m3=matrix(v2,nrow=2,ncol=4, byrow=T)
> m3
     [,1] [,2] [,3] [,4]
[1,]    3    4    3    4
[2,]    3    4    3    4
```

Lists

- Collection of objects

```
> d<-15;
> a<-matrix(c(1,2,3,4),2,2);
> a
      [,1] [,2]
[1,]    1    3
[2,]    2    4
> b<-list(first=d, second=a, x="mary")
> b
$first
[1] 15

$second
      [,1] [,2]
[1,]    1    3
[2,]    2    4
> b$x
[1] "mary"
> b[[3]]
[1] "mary"
```

Factors

- Categorize data, e.g.
→ Important for classification

```
> f1=c("Man", "Woman")
> f1
[1] "Man"    "Woman"
> f2=c("Man", "Woman", "Man")
> table(f2)
f2
   Man Woman
      2     1
> f3=factor(c(1,0,1,1,0), levels=c(0,1), labels=c("Man", "Woman"))
> f3
[1] Woman Man    Woman Woman Man
Levels: Man Woman
>
```

Data frame

- Vectors and matrices of the same row length can be collected into a data frame
 - Used to store data of different types from same observation into a single table
 - Use: `data.frame(object 1, object 2, ... , object k)`

```
> x<-c(1,3)
> y<-c("M", "F")
> z<-data.frame(x,y)
> z
  x y
1 1 M
2 3 F
```

Data frame

- Any column in the data frame can be retrieved by
dataframe\$object or dataframe[[col_nr]]

```
> z$x
[1] 1 3
> z[[1]]
[1] 1 3
> z$y
[1] M F
Levels: F M
```

- Any row in the data frame can be extracted by using matrix notation, e.g.,
z [1,]

Read data from CSV files

- 1. Specify exact file path or change current directory → Set Working Directory or `setwd()`
- 2. Use:
 - `dataframe <- read.csv2(file_name)`
→ comma as decimals, semi-colon as separator
 - `dataframe <- read.csv(file_name)`
→ dot as decimals, comma as separator
- Similar workflow for other file types , e.g. read from Excel files with packages `xlsx`, `readxl`

Conversion between data types

- E.g.,

- Data frame to matrix
- Matrix to data frame
- Numeric to factor
- Factor to numeric
- List to vector
- Vector to list

```
> v6=c(1,4,2,2,1)
> f4=as.factor(v6)
> f4
[1] 1 4 2 2 1
Levels: 1 2 4
> f5=c("1", "0", "1", "1", "1")
> f5
[1] "1" "0" "1" "1" "1"
> as.list(f5)
[[1]]
[1] "1"

[[2]]
[1] "0"

[[3]]
[1] "1"

[[4]]
[1] "1"

[[5]]
[1] "1"

> l1=list(a=1, b=3)
> l1
$a
[1] 1
$b
[1] 3

> as.numeric(l1)
[1] 1 3
```

Loops

```
for (name in expr1 )  
{  
...  
}  
  
while (condition)  
{  
...  
}
```

```
> for (i in 1:5) {  
+ y=seq(i,8)  
+ print(y) }  
[1] 1 2 3 4 5 6 7 8  
[1] 2 3 4 5 6 7 8  
[1] 3 4 5 6 7 8  
[1] 4 5 6 7 8  
[1] 5 6 7 8  
> |
```

Loops with conditions

```
if  (x==3)  {  
...  
...  
} else {  
...  
}  
  
> m4=matrix(c(1,2,0,1), nrow=2)  
> m4  
[1,] [,1] [,2]  
[1,]    1    0  
[2,]    2    1  
> n=dim(m4)[1]  
> I=numeric(n)  
> for (i in 1:n){  
+   if(max(m4[i,]>1)) I[i]=1  
+ }  
> I  
[1] 0 1
```

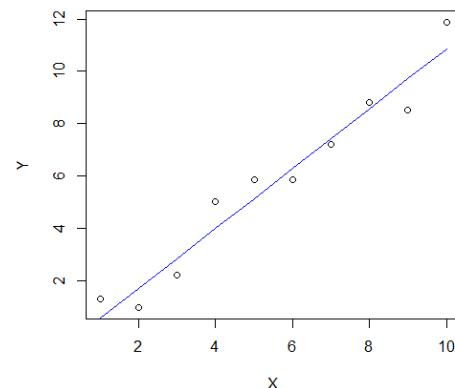
Using functions

→ We have been using many functions already throughout this lecture

- There are some obligatory parameters and optional parameters
 - The optional parameters can be specified in different order
- Use `?name_of_function` to see function parameters, e.g. `?lm`

```
X=1:10  
Y=1:10+rnorm(10)  
W=c(rep(1,5), rep(2,5))  
mydata=data.frame(X,Y)  
  
result=lm(Y~X, weights=W,data=mydata)  
?predict.lm  
Fit=predict(result)
```

```
plot(X,Y)  
points(X,Fit, type="l", col="blue")
```



Random number generation

- Random is not random
→ **Use `set.seed(12345)` to get identical results**
- Plenty of random number generators available
 - `Rnorm`
 - `Runif`
 - ...
- Use `d` for density, `p` for CDF, `q` for quantiles, and `r` for simulation: (ex: `rnorm pnorm dnorm qnorm`)

Distribution	R name	additional arguments
beta	beta	<code>shape1, shape2, ncp</code>
binomial	binom	<code>size, prob</code>
Cauchy	cauchy	<code>location, scale</code>
chi-squared	chisq	<code>df, ncp</code>
exponential	exp	<code>rate</code>
F	f	<code>df1, df2, ncp</code>
gamma	gamma	<code>shape, scale</code>
geometric	geom	<code>prob</code>
hypergeometric	hyper	<code>m, n, k</code>
log-normal	lnorm	<code>meanlog, sdlog</code>
logistic	logis	<code>location, scale</code>
negative binomial	nbinom	<code>size, prob</code>
normal	norm	<code>mean, sd</code>
Poisson	pois	<code>lambda</code>
Student's t	t	<code>df, ncp</code>
uniform	unif	<code>min, max</code>
Weibull	weibull	<code>shape, scale</code>
Wilcoxon	wilcox	<code>m, n</code>

Writing your own functions

- Define a function with `x <- function() {}`
→ {} define function scope, i.e., what the function will do
- Specify which parameters a functions takes
- Specify what values a functions returns with `return(value)`

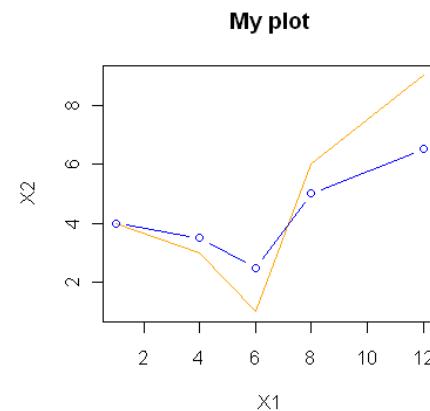
```
> myfun <-function(x=20, y, z)
+ {
+ if(z)
+ result=x+y
+ else
+ {
+ result=log(x)*y
+ }
+ result
+ }
> r<-myfun(1,2, TRUE)
> r
[1] 3
> r<-myfun(z=FALSE, y=0)
> r
[1] 0
> |
```

Graphical procedures

- Some common procedures:
 - `plot(x,..)` plots time series
 - `plot(x,y)` scatter plot
 - `plot(x,y)` followed by `points(x,y)` plots several scatterplots in one coordinate system
 - `hist(x,..)` plots a histogram
 - `persp(x,y,z,...)` creates surface plots
 - `cloud(formula,data..)` creates 3D scatter plot

```
x<-c(1,4,7,8,12);
y<-c(4,3,1,6,9);

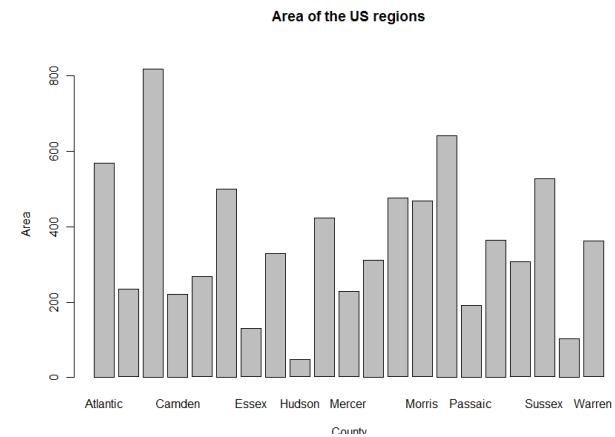
plot(x,y, type="l", col="orange",
      main="My plot", xlab="X1", ylab="X2");
points(x, y/2+2, type="b", col="blue");
```



Graphical procedures

- Parameters:
 - Color: col="color"
 - Title: main="text"
 - Footnote: sub="text"
 - X-axis label: xlab="text"
 - Y-axis label: ylab="text"

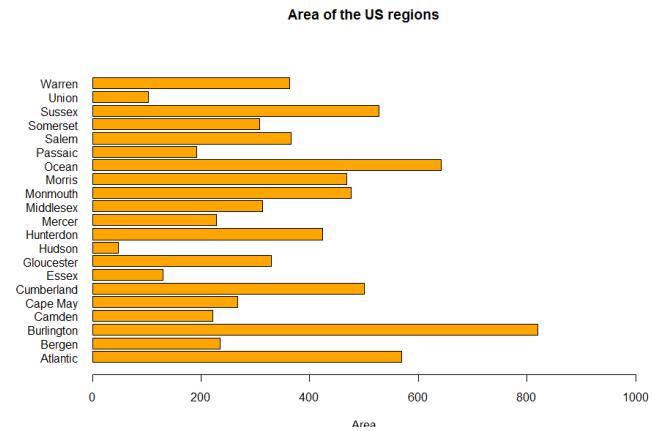
```
mydata<-read.csv2("Counties.csv");
barplot(mydata$Area,
names.arg=mydata$County, main="Area of
the US regions", xlab="County",
ylab="Area");
```



Graphical procedures

- Some parameters need to be specified either in the plotting function or inside `par(...)`
 - Symbol that is plotted: `pch=number`
 - Linetype: `lty=number`
 - Direction of axis values: `las=0, 1 or 2`
 - Margins (inch): `mai=c(bottom, left, top, right)`
 - Horizontal justification: `adj=between 0 and 1`

```
barplot(mydata$Area,  
names.arg=mydata$County,  
horiz=TRUE, las=1,  
xlim=c(0,1000), col="orange",  
main="Area of the US  
regions", xlab="Area");
```



Useful code snippets

- Dividing a dataset into training/test sets:

```
data=data.frame(X=c(1,1,2,2,3), Y=c("M","F","M","M","F"))
n=dim(data)[1]
set.seed(12345)
id=sample(1:n, floor(n*0.5))
train=data[id,]
test=data[-id,]
```

```
> train
  X Y
4 2 M
5 3 F
> test
  X Y
1 1 M
2 1 F
3 2 M
.
```

- Computing misclassification rate

```
missclass=function(X,X1){
  n=length(X)
  return(1-sum(diag(table(X,X1)))/n)
}
```

```
> X=c(1,3,1,1,2,3,1,2,2,2,1,1,3)
> Xfit=c(2,3,2,1,2,3,1,2,2,2,1,1,1)
> missclass(X,Xfit)
[1] 0.2307692
```

Additional resources

- R manuals of dplyr and tidyr packages (data manipulation) → many useful functions to apply on datasets
- R language manual <https://cran.r-project.org/doc/manuals/r-release/R-intro.pdf>